

# Load the dataset

```
In [3]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline


from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, f1_score
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from imblearn.over_sampling import SMOTE
from sklearn.linear_model import LogisticRegression
```

```
In [5]: df = pd.read_csv('creditcard.csv')
df.head()
```

```
Out[5]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278

5 rows × 31 columns



```
In [6]: # statistical info
df.describe()
```

```
Out[6]:
```

	Time	V1	V2	V3	V4	V5	V6	V7
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15	9.604066e-16	1.487313e-15	-5.556467e-16
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380247e+00	1.332271e+00	1.237094e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+02	-2.616051e+01	-4.355724e+01
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915971e-01	-7.682956e-01	-5.540759e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433583e-02	-2.741871e-01	4.010308e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119264e-01	3.985649e-01	5.704361e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480167e+01	7.330163e+01	1.205895e+02

8 rows × 31 columns



```
In [11]: # datatype info
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Time        284807 non-null  float64
 1   V1          284807 non-null  float64
 2   V2          284807 non-null  float64
 3   V3          284807 non-null  float64
 4   V4          284807 non-null  float64
 5   V5          284807 non-null  float64
 6   V6          284807 non-null  float64
 7   V7          284807 non-null  float64
 8   V8          284807 non-null  float64
 9   V9          284807 non-null  float64
10  V10         284807 non-null  float64
11  V11         284807 non-null  float64
12  V12         284807 non-null  float64
13  V13         284807 non-null  float64
14  V14         284807 non-null  float64
15  V15         284807 non-null  float64
16  V16         284807 non-null  float64
17  V17         284807 non-null  float64
18  V18         284807 non-null  float64
19  V19         284807 non-null  float64
20  V20         284807 non-null  float64
21  V21         284807 non-null  float64
22  V22         284807 non-null  float64
23  V23         284807 non-null  float64
24  V24         284807 non-null  float64
25  V25         284807 non-null  float64
26  V26         284807 non-null  float64
27  V27         284807 non-null  float64
28  V28         284807 non-null  float64
29  Amount      284807 non-null  float64
30  Class       284807 non-null  int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB

```

## Preprocessing the dataset

```

In [13]: # check for null values
df.isnull().sum()

```

```


Out[13]: Time        0
V1            0
V2            0
V3            0
V4            0
V5            0
V6            0
V7            0
V8            0
V9            0
V10           0
V11           0
V12           0
V13           0
V14           0
V15           0
V16           0
V17           0
V18           0
V19           0
V20           0
V21           0
V22           0
V23           0
V24           0
V25           0
V26           0
V27           0
V28           0
Amount        0
Class         0
dtype: int64

```

## Exploratory Data Analysis

```
In [15]: sns.countplot(x='Class', data=df)
```

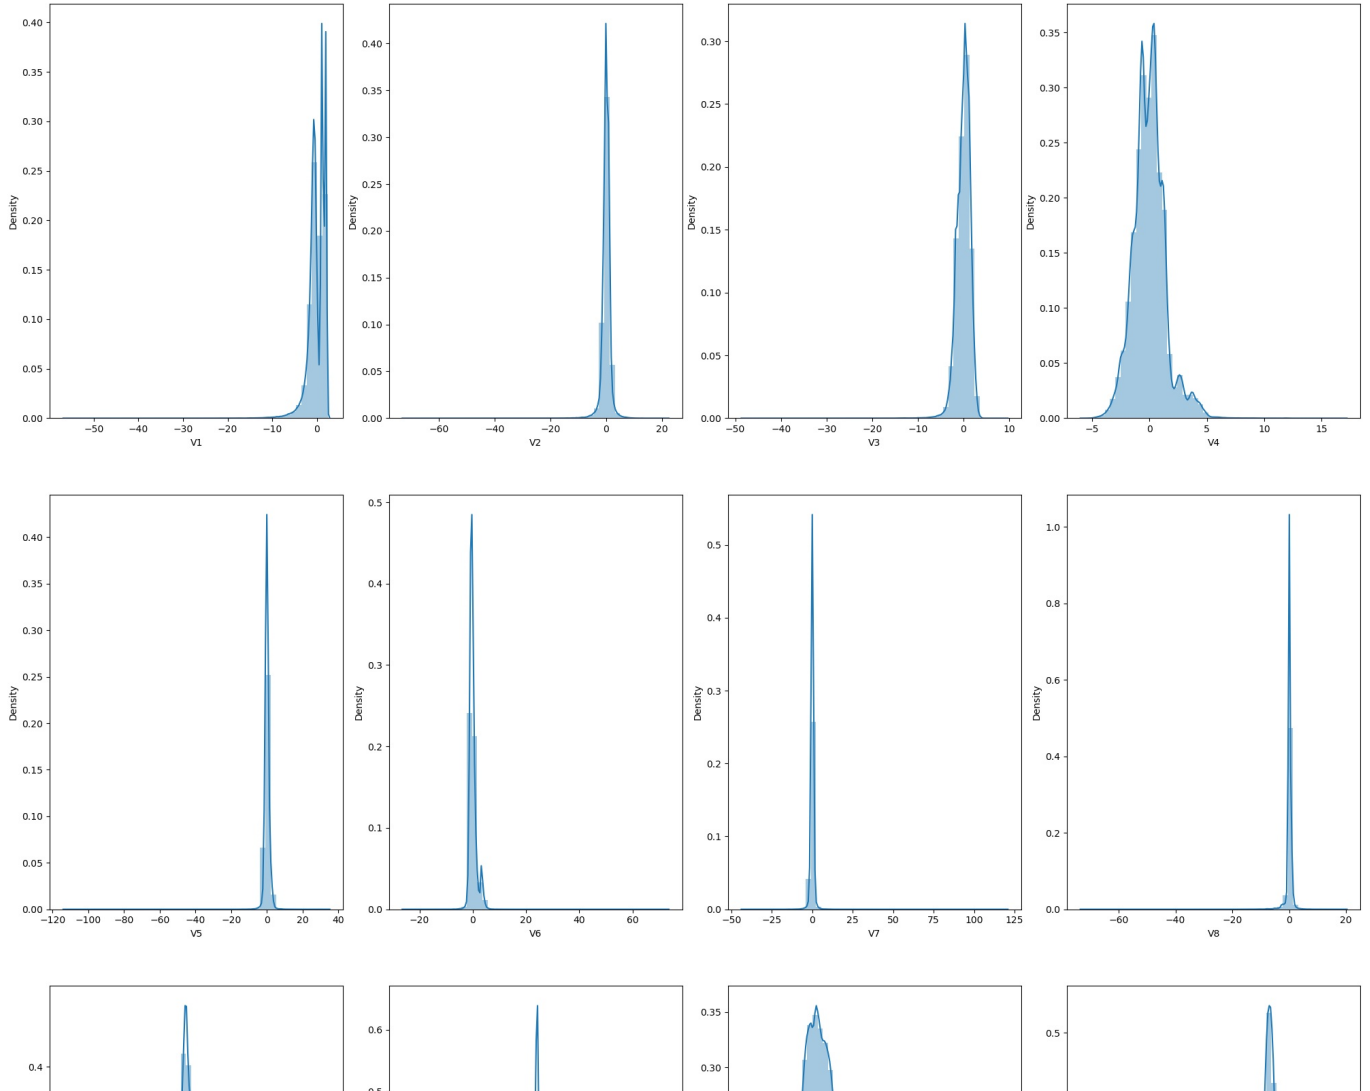
```
Out[15]: <Axes: xlabel='Class', ylabel='count'>
```

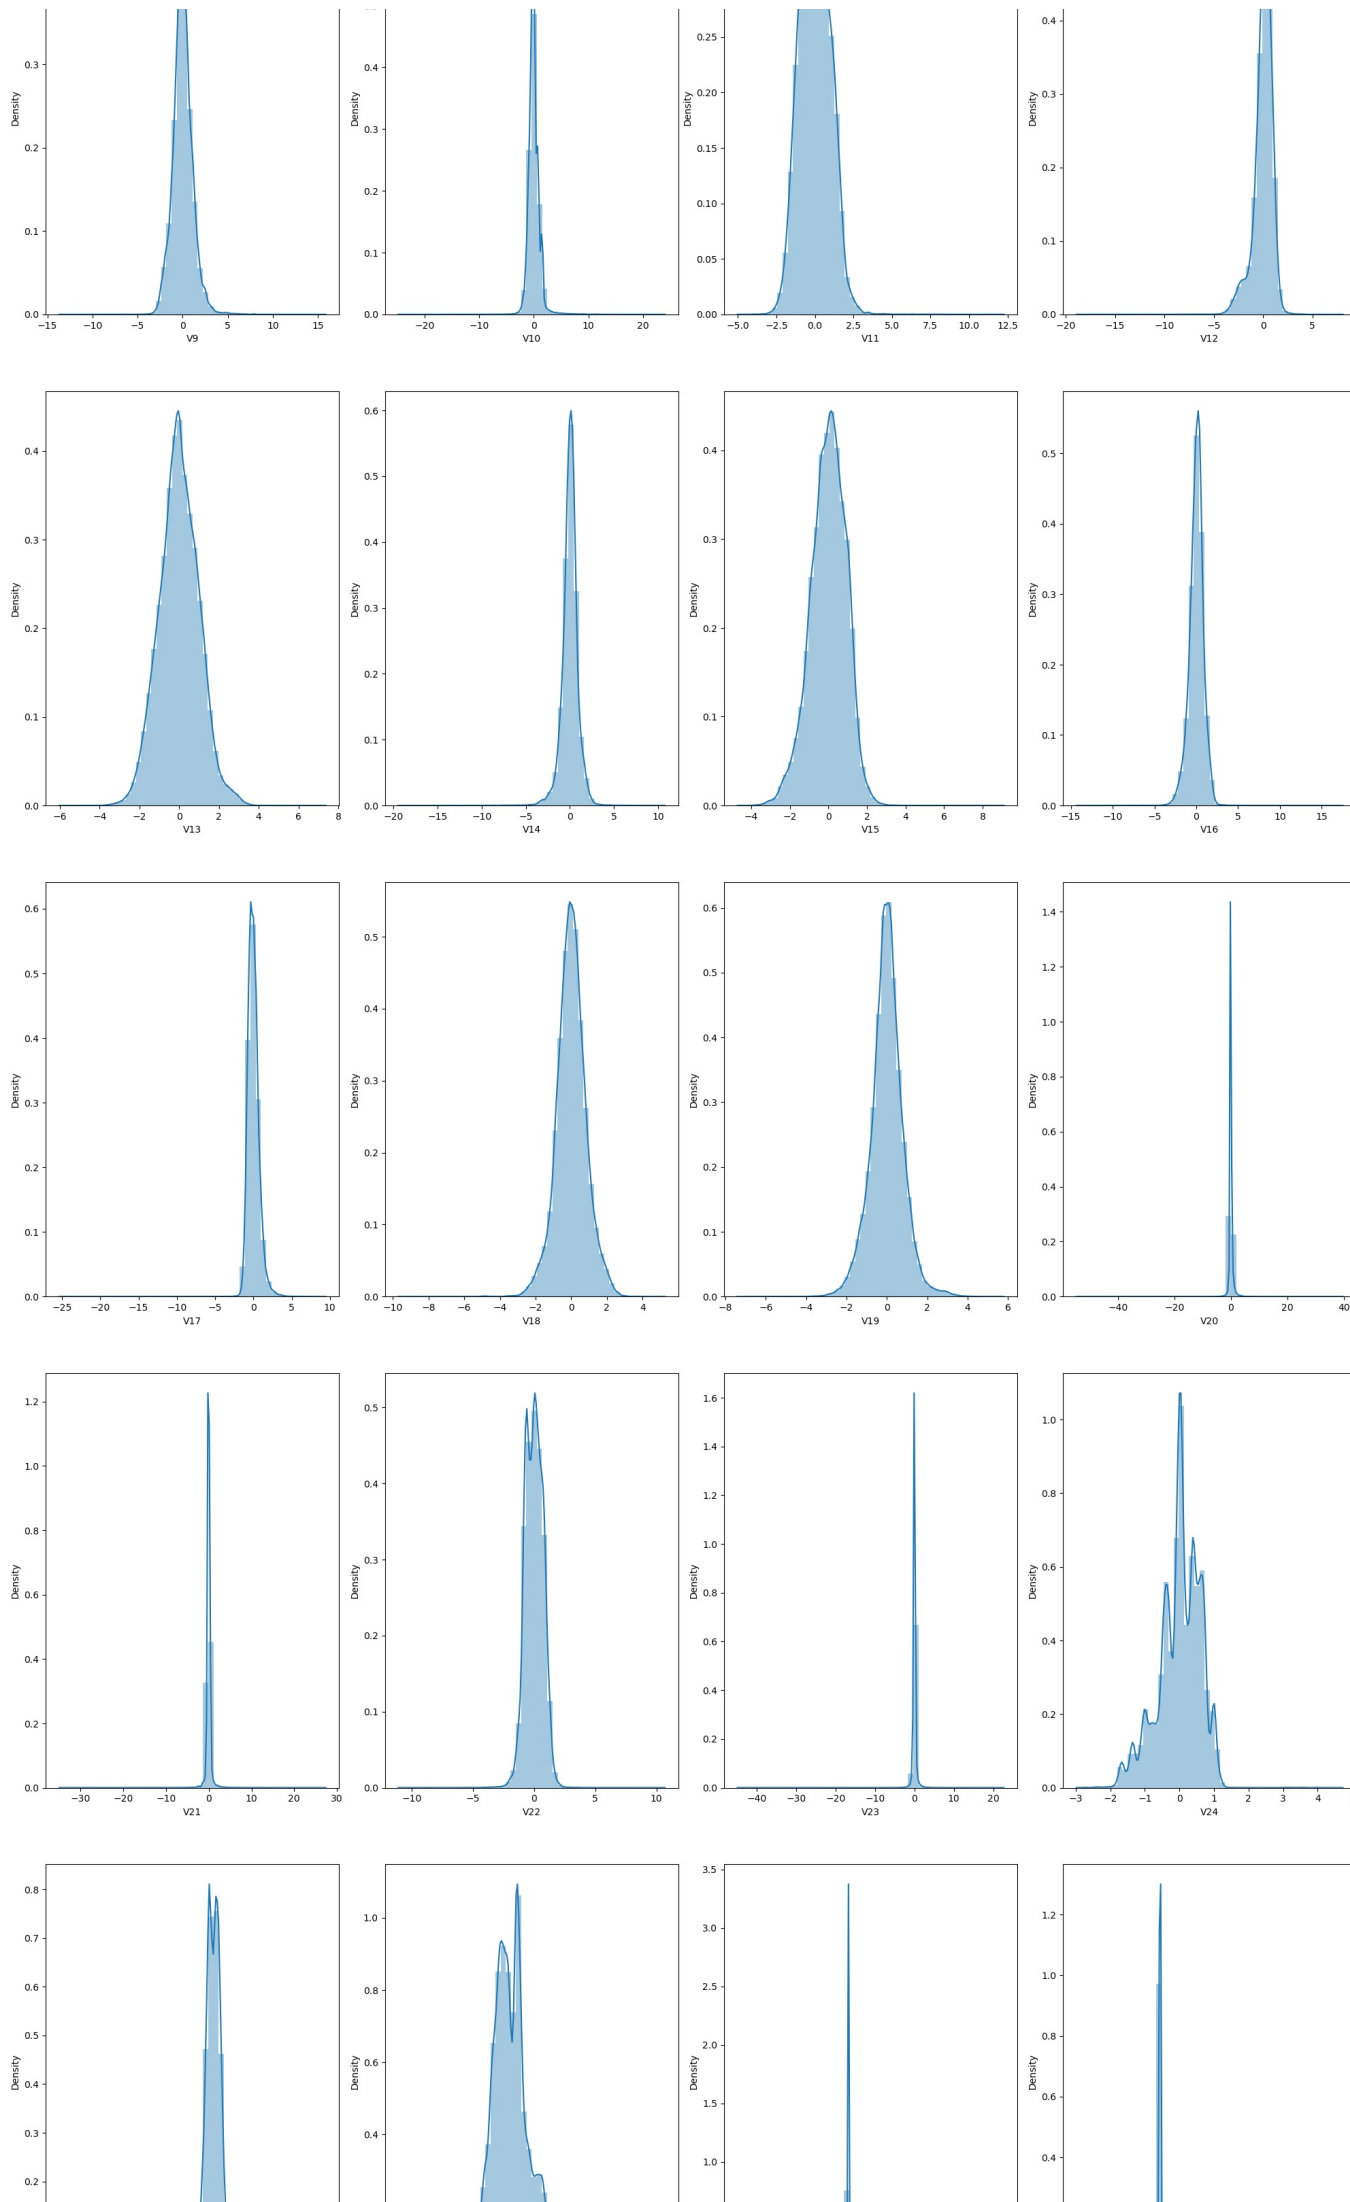



```
In [17]: df_temp = df.drop(columns=['Time', 'Amount', 'Class'], axis=1)
```

```
# create dist plots
fig, ax = plt.subplots(ncols=4, nrows=7, figsize=(20, 50))
index = 0
ax = ax.flatten()

for col in df_temp.columns:
    sns.distplot(df_temp[col], ax=ax[index])
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.5, h_pad=5)
```









```
In [18]: sns.distplot(df['Time'])
```

```
Out[18]: <Axes: xlabel='Time', ylabel='Density'>
```



```
In [19]: sns.distplot(df['Amount'])
```


```
Out[19]: <Axes: xlabel='Amount', ylabel='Density'>
```



## Coorelation Matrix

```
In [20]: corr = df.corr()
plt.figure(figsize=(20,20))
sns.heatmap(corr, annot=True, cmap='coolwarm')
```

```
Out[20]: <Axes: >
```



## Input Split

```
In [25]: X = df.drop(columns=['Class'], axis=1)
         y = df['Class']
```

## Standard Scaling

```
In [27]: sc = StandardScaler()
         x_scaler = sc.fit_transform(X)
```

```
In [29]: x_scaler[-1]
```

```
Out[29]: array([ 1.64205773, -0.27233093, -0.11489898,  0.46386564, -0.35757
,
-0.00908946, -0.48760183,  1.27476937, -0.3471764 ,  0.44253246,
-0.84072963, -1.01934641, -0.0315383 , -0.18898634, -0.08795849,
 0.04515766, -0.34535763, -0.77752147,  0.1997554 , -0.31462479,
 0.49673933,  0.35541083,  0.8861488 ,  0.6033653 ,  0.01452561,
-0.90863123, -1.69685342, -0.00598394,  0.04134999,  0.51435531])
```

# Model Training

```
In [31]: x_train, x_test, y_train, y_test = train_test_split(x_scaler, y, test_size=0.25, random_state=42, stratify=y)
```

```
In [33]: model = LogisticRegression()
# training
model.fit(x_train, y_train)
# testing
y_pred = model.predict(x_test)
print(classification_report(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	71079
1	0.84	0.62	0.71	123
accuracy			1.00	71202
macro avg	0.92	0.81	0.85	71202
weighted avg	1.00	1.00	1.00	71202

F1 Score: 0.7102803738317757

```
In [35]: model = RandomForestClassifier()
# training
model.fit(x_train, y_train)
# testing
y_pred = model.predict(x_test)
print(classification_report(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	71079
1	0.95	0.79	0.86	123
accuracy			1.00	71202
macro avg	0.98	0.89	0.93	71202
weighted avg	1.00	1.00	1.00	71202

F1 Score: 0.8622222222222222

```
In [39]: model = XGBClassifier(n_jobs=-1)
# training
model.fit(x_train, y_train)
# testing
y_pred = model.predict(x_test)
print(classification_report(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
```


	precision	recall	f1-score	support
0	1.00	1.00	1.00	71079
1	0.95	0.79	0.86	123
accuracy			1.00	71202
macro avg	0.98	0.89	0.93	71202
weighted avg	1.00	1.00	1.00	71202

F1 Score: 0.8622222222222222

# class Imbalancement

```
In [41]: sns.countplot(x=y_train, data=df)
```


```
Out[41]: <Axes: xlabel='Class', ylabel='count'>
```



```
In [43]: over_sample = SMOTE()  
x_smote, y_smote = over_sample.fit_resample(x_train, y_train)
```

```
In [45]: sns.countplot(x=y_smote, data=df)
```

```
Out[45]: <Axes: xlabel='Class', ylabel='count'>
```



```
In [47]: model = LogisticRegression()  
# training  
model.fit(x_smote, y_smote)  
# testing  
y_pred = model.predict(x_test)  
print(classification_report(y_test, y_pred))  
print("F1 Score:", f1_score(y_test, y_pred))
```



	precision	recall	f1-score	support
0	1.00	0.98	0.99	71079
1	0.06	0.89	0.11	123
accuracy			0.98	71202
macro avg	0.53	0.93	0.55	71202
weighted avg	1.00	0.98	0.99	71202

F1 Score: 0.11139499233520694

```
In [49]: model = RandomForestClassifier(n_jobs=-1)
# training
model.fit(x_smote, y_smote)
# testing
y_pred = model.predict(x_test)
print(classification_report(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	71079
1	0.89	0.80	0.84	123
accuracy			1.00	71202
macro avg	0.95	0.90	0.92	71202
weighted avg	1.00	1.00	1.00	71202

F1 Score: 0.8412017167381974

```
In [51]: model = XGBClassifier(n_jobs=-1)
# training
model.fit(x_smote, y_smote)
# testing
y_pred = model.predict(x_test)
print(classification_report(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	71079
1	0.75	0.84	0.79	123
accuracy			1.00	71202
macro avg	0.88	0.92	0.90	71202
weighted avg	1.00	1.00	1.00	71202

F1 Score: 0.7923076923076923

```
In [53]: # created by Piyush Verma as Data Scientist at Growintern
# For contact pverma20269968@gmail.com
```

```
In [ ]:
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js